

iTRANSFORMER: INVERTED TRANSFORMERS ARE EFFECTIVE FOR TIME SERIES FORECASTING

Yong Liu*, **Tengge Hu***, **Haoran Zhang***, **Haixu Wu**, **Shiyu Wang[§]**, **Lintao Ma[§]**, **Mingsheng Long[✉]**

School of Software, BNRist, Tsinghua University, Beijing 100084, China

[§]Ant Group, Hangzhou, China

{liuyong21, htg21, z-hr20, whx20}@mails.tsinghua.edu.cn

{weiming.wsy, lintao.mlt}@antgroup.com, mingsheng@tsinghua.edu.cn

ABSTRACT

The recent boom of linear forecasting models questions the ongoing passion for architectural modifications of Transformer-based forecasters. These forecasters leverage Transformers to model the global dependencies over *temporal tokens* of time series, with each token formed by multiple variates of the same timestamp. However, **Transformer is challenged in forecasting series with larger lookback windows due to performance degradation and computation explosion.** Besides, the unified embedding for each temporal token fuses multiple variates with potentially unaligned timestamps and distinct physical measurements, which may fail in learning variate-centric representations and result in meaningless attention maps. In this work, we reflect on the competent duties of Transformer components and repurpose the Transformer architecture without any adaptation on the basic components. We propose **iTransformer** that simply inverts the duties of the attention mechanism and the feed-forward network. Specifically, the time points of individual series are embedded into *variate tokens* which are utilized by the attention mechanism to capture multivariate correlations; meanwhile, the feed-forward network is applied for each variate token to learn nonlinear representations. The iTransformer model achieves consistent state-of-the-art on several real-world datasets, which further empowers the Transformer family with promoted performance, generalization ability across different variates, and better utilization of arbitrary lookback windows, making it a nice alternative as the fundamental backbone of time series forecasting.

1 INTRODUCTION

Transformer (Vaswani et al., 2017) has achieved tremendous success in natural language processing (Brown et al., 2020) and computer vision (Dosovitskiy et al., 2021), growing into the foundation model that follows the scaling law (Kaplan et al., 2020). Inspired by the immense success in extensive fields, Transformer with strong capabilities of depicting pairwise dependencies and extracting multi-level representations in sequences is emerging in time series forecasting (Li et al., 2021; Wu et al., 2021; Nie et al., 2023).

However, researchers have recently begun to question the validity of Transformer-based forecasters, which typically embed multiple variates of the same timestamp into indistinguishable channels and apply attention on these *temporal tokens* to capture temporal dependencies. Considering the numerical but less semantic relationship among time points, researchers find that simple linear layers, which can be traced back to statistical forecasters (Box & Jenkins, 1968), have exceeded complicated Transformers on both performance and efficiency (Zeng et al., 2023; Das et al., 2023). Meanwhile, ensuring the independence of variate

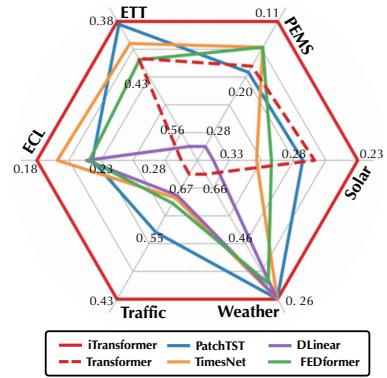


Figure 1: Performance of iTransformer. Average results (MSE) are reported following TimesNet (2023).

*Equal Contribution

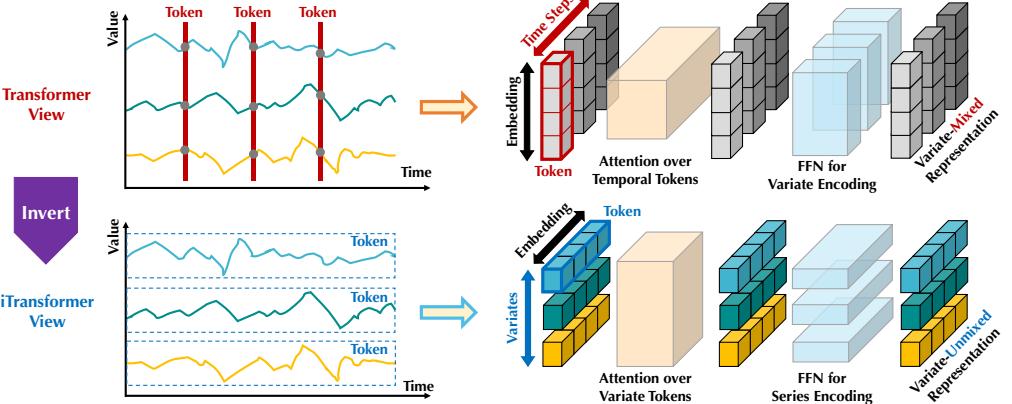


Figure 2: Comparison between the vanilla Transformer (top) and the proposed iTransformer (bottom). Unlike Transformer, which embeds each time step to the temporal token, iTransformer embeds the whole series independently to the variate token, such that multivariate correlations can be depicted by the attention mechanism and series representations are encoded by the feed-forward network.

and utilizing mutual information is ever more highlighted by recent research that explicitly models multivariate correlations to achieve accurate forecasting (Zhang & Yan, 2023; Ekambaram et al., 2023), but this goal can be hardly achieved without subverting the vanilla Transformer architecture.

Considering the disputes of Transformer-based forecasters, we reflect on why Transformers perform even worse than linear models in time series forecasting while acting predominantly in many other fields. We notice that the existing structure of Transformer-based forecasters may be not suitable for multivariate time series forecasting. As shown in the left of Figure 2, it is notable that the points of the same time step that basically represent completely different physical meanings recorded by inconsistent measurements are embedded into one token with wiped-out multivariate correlations. And the token formed by a single time step can struggle to reveal beneficial information due to excessively local receptive field and unaligned timestamps of multivariate time points in the real world. Besides, while series variations can be greatly influenced by the sequence order, permutation-invariant attention mechanisms are improperly adopted on the temporal dimension (Zeng et al., 2023). Consequently, Transformer is weakened to capture essential series representations and portray multivariate correlations, limiting its capacity and generalization ability on diverse time series data.

Concerning the irrationality of embedding multivariate points of each time step as a (temporal) token, we take an *inverted view* on time series and embed the whole time series of each variate independently into a (variate) token, the extreme case of Patching (Nie et al., 2023) that enlarges local receptive field. By inverting, the embedded token aggregates the global representations of series that can be more variate-centric and better leveraged by booming attention mechanisms for multivariate correlating. Meanwhile, the feed-forward network can be proficient enough to learn generalizable representations for distinct variates encoded from arbitrary lookback series and decoded to predict future series.

Based on the above motivations, we believe it is not that Transformer is ineffective for time series forecasting, but rather it is improperly used. In this paper, we revisit the structure of Transformer and advocate *iTransformer* as a fundamental backbone for time series forecasting. Technically, we embed each time series as *variate tokens*, adopt the attention for multivariate correlations, and employ the feed-forward network for series encoding. Experimentally, the proposed iTransformer achieves state-of-the-art performance on real-world forecasting benchmarks shown in Figure 1 and surprisingly tackles the pain points of Transformer-based forecasters. Our contributions lie in three aspects:

- We reflect on the architecture of Transformer and refine that the competent capability of native Transformer components on time series is underexplored.
- We propose iTransformer that regards independent time series as tokens to capture multivariate correlations by self-attention and utilize layer normalization and feed-forward network modules to learn better series-global representations for time series forecasting.
- Experimentally, iTransformer achieves consistent state-of-the-art on real-world forecasting benchmarks. We extensively analyze the inverted modules and architecture choices, indicating a promising direction for the future improvement of Transformer-based forecasters.

2 RELATED WORK

With the progressive breakthrough made in natural language processing and computer vision areas, elaboratively designed Transformer variants are proposed to tackle ubiquitous time series forecasting applications. Going beyond contemporaneous TCNs (Bai et al., 2018; Liu et al., 2022a) and RNN-based forecasters (Zhao et al., 2017; Rangapuram et al., 2018; Salinas et al., 2020), Transformer has exhibited powerful sequence modeling capability and promising model scalability, leading to the trend of passionate modifications adapted for time series forecasting.

Through a systematical review of Transformer-based forecasters, we conclude that existing modifications can be divided into four categories by whether to modify the component and architecture. As shown in Figure 3, the first category (Wu et al., 2021; Li et al., 2021; Zhou et al., 2022), which is the most common practice, mainly concerns the component adaptation, especially the attention module for the temporal dependency modeling and the complexity optimization on long sequences. Nevertheless, with the rapid emergence of linear forecasters (Oreshkin et al., 2019; Zeng et al., 2023; Das et al., 2023), the impressive performance and efficiency continuously challenge this direction. Soon afterward, the second category attempts to fully utilize Transformer and pays more attention to inherent processing of time series, such as Stationarization (Liu et al., 2022b), Channel Independence, and Patching (Nie et al., 2023), which bring about consistently improved performance. Moreover, faced with the increasing significance of the independence and mutual interactions of multiple variates, the third category refurbishes Transformer in both aspects of component and architecture. Representative (Zhang & Yan, 2023) explicitly captures the cross-time and cross-variate dependency by the renovated attention mechanism and architecture.

Unlike previous works, our proposed model modifies none of the native components of Transformer. By contrast, the duties of components are totally inverted, which is the only one that belongs to the fourth category to our best knowledge. We believe the capabilities of the original modules have stood the test by extensive fields, the truth is that the architecture of Transformer is improperly adopted.

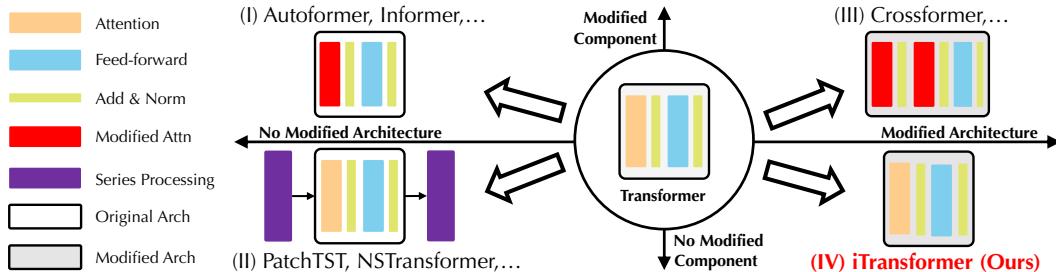


Figure 3: Transformer-based forecasters categorized by component and architecture modifications.

3 iTRANSFORMER

In multivariate time series forecasting, given historical observations $\mathbf{X} = \{\mathbf{x}_1, \dots, \mathbf{x}_T\} \in \mathbb{R}^{T \times N}$ with T time steps and N variates, we predict the future S time steps $\mathbf{Y} = \{\mathbf{x}_{T+1}, \dots, \mathbf{x}_{T+S}\} \in \mathbb{R}^{S \times N}$. For convenience, we denote $\mathbf{X}_{t,:}$ as the simultaneously recorded multivariate at time step t , and $\mathbf{X}_{:,n}$ as the whole time series of each variate indexed by n . It is notable that $\mathbf{X}_{t,:}$ may not contain time points of the essentially same timestamp in real-world scenarios because of the systematical delay of monitors and loosely organized datasets. The elements of $\mathbf{X}_{t,:}$ can be distinct from each other in physical measurements and statistical distributions, for which a variate $\mathbf{X}_{:,n}$ generally shares.

3.1 STRUCTURE OVERVIEW

Our proposed *iTransformer* illustrated in Figure 4 utilizes the simpler encoder-only architecture of Transformer (Vaswani et al., 2017), including the embedding, projection, and Transformer blocks.

Embedding the whole series as the token Most Transformer-based forecasters typically regard multiple variates of the same time as the (temporal) token and follow the generative formulation of forecasting tasks. However, we find the approach on the numerical modality can be less instructive for

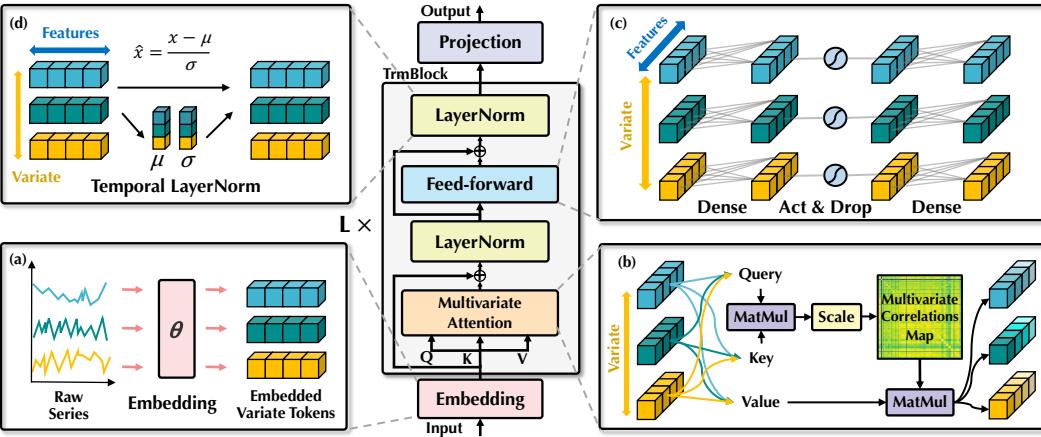


Figure 4: Overall structure of iTransformer, which shares the same modular arrangement with the encoder of Transformer: (a) raw series of different variates are independently embedded to tokens. (b) self-attention is applied to embedded variate tokens with enhanced interpretability revealing multivariate correlations. (c) series representations of each token are extracted by the shared feed-forward network. (d) layer normalization is adopted to reduce the discrepancies among variates.

learning attention maps, which is supported by increasing applications of Patching (Dosovitskiy et al., 2021; Nie et al., 2023) that broadens the respective field. Meanwhile, the triumph of linear forecasters also challenges the necessity of adopting a heavy encoder-decoder Transformer for generating tokens. Instead, our proposed encoder-only iTransformer focuses on representation learning and adaptive correlating of multivariate series. Each time series driven by the underlying complicated process is firstly tokenized to describe the properties of the variate, applied by self-attention for mutual interactions, and individually processed by feed-forward networks for series representations. Notably, the task to generate the predicted series is essentially delivered to linear layers, which has been proven competent by previous work (Das et al., 2023) and we provide a detailed analysis in the next section.

Based on the above considerations, in iTransformer, the process of predicting future series of each specific variate $\hat{Y}_{:,n}$ based on the lookback series $\mathbf{X}_{:,n}$ is simply formulated as follows:

$$\begin{aligned} \mathbf{h}_n^0 &= \text{Embedding}(\mathbf{X}_{:,n}), \\ \mathbf{H}^{l+1} &= \text{TrmBlock}(\mathbf{H}^l), \quad l = 0, \dots, L-1, \\ \hat{\mathbf{Y}}_{:,n} &= \text{Projection}(\mathbf{h}_n^L), \end{aligned} \quad (1)$$

where $\mathbf{H} = \{\mathbf{h}_1, \dots, \mathbf{h}_N\} \in \mathbb{R}^{N \times D}$ contains N embedded tokens of dimension D and the superscript denotes the layer index. $\text{Embedding} : \mathbb{R}^T \mapsto \mathbb{R}^D$ and $\text{Projection} : \mathbb{R}^D \mapsto \mathbb{R}^S$ are both implemented by multi-layer perceptron (MLP). The obtained variate tokens interact with each other by self-attention and are independently processed by the shared feed-forward network in each TrmBlock. Specifically, as the order of sequence is implicitly stored in the neuron permutation of the feed-forward network, the position embedding in the vanilla Transformer is no longer needed here.

iTransformers The proposed architecture essentially presupposes no more specific requirements on Transformer variants, other than the attention mechanism should be applicable for multivariate correlation modeling. Thus, a bundle of efficient attention mechanisms (Li et al., 2021; Wu et al., 2022; Dao et al., 2022) can be the plugins, reducing the complexity of correlating when the variate number grows large. The equipped Transformer variants by the proposed architecture, which we name as *iTransformers*, are extensively evaluated in our experiments in Section 4.2 and demonstrate superior advantages on time series forecasting.

3.2 INVERTED TRANSFORMER COMPONENTS

We organize a stack of L blocks composed of native layer normalization, feed-forward network, and self-attention modules. But their duties on the inverted dimension should be carefully reconsidered.

Layer normalization Layer normalization (Ba et al., 2016) is originally proposed to increase the convergence and training stability of deep networks. In typical Transformer-based forecasters, the module normalizes the variate representation of the same timestamp, gradually making each variate indistinguishable from each other. Once the collected time points are not aligned by time, the operation will also introduce interaction noises between noncausal or delayed processes. In our inverted version, the normalization is applied to the series representation of individual variate as Equation 2, which has been studied and proved effective in tackling non-stationary problems (Kim et al., 2021; Liu et al., 2022b). Besides, as all series as (variante) tokens are normalized to a normal distribution, the discrepancies caused by inconsistent measurements can be diminished. Instead, in previous architecture, different tokens of time steps will be normalized, leading to oversmooth series.

$$\text{LayerNorm}(\mathbf{H}) = \left\{ \frac{\mathbf{h}_n - \text{Mean}(\mathbf{h}_n)}{\sqrt{\text{Var}(\mathbf{h}_n)}} \mid n = 1, \dots, N \right\} \quad (2)$$

Feed-forward network Transformer leverages the feed-forward network as the basic building block for encoding token representation and it is identically applied to each token. As aforementioned, in the vanilla Transformer, multiple variates of the same timestamp that forms the token can be malpositioned and too localized to reveal enough information for predictions. In the inverted version, feed-forward networks are leveraged on the channels of variate tokens. By the universal approximation theorem (Hornik, 1991), they can extract complicated representations to describe a time series. With the stacking of inverted blocks, they are devoted to encoding the observed time series and decoding the representations for future series using dense non-linear connections, which works the same as the recent works completely built on MLPs (Tolstikhin et al., 2021; Das et al., 2023).

More interestingly, the identical linear operation on independent time series, which serves as the combination of the recent linear forecasters (Zeng et al., 2023) and Channel Independent strategy (Nie et al., 2023), can be instructive for us to understand the series representations. Recent revisiting on linear forecasters (Li et al., 2023) highlights that temporal features extracted by MLPs are supposed to be shared within distinct time series. We propose a rational explanation that the neurons of MLP are taught to portray the intrinsic properties of any time series, such as the amplitude, periodicity, and even frequency spectrums (neuron as a filter), serving as a more advantageous predictive representation learner than the self-attention applied on time points. Experimentally, we validate that the division of labor helps enjoy the benefits of linear layers in Section 4.3, such as the promoted performance if providing enlarged lookback series, and the generalization ability on unseen variates.

Self-attention While the attention mechanism is generally adopted for facilitating the temporal dependencies modeling in previous forecasters, the inverted model regards the whole series of one variate as an individual process. Concretely, with comprehensively extracted representations of each time series $\mathbf{H} = \{\mathbf{h}_0, \dots, \mathbf{h}_N\} \in \mathbb{R}^{N \times D}$, the self-attention module adopts linear projections to get queries, keys and values $\mathbf{Q}, \mathbf{K}, \mathbf{V} \in \mathbb{R}^{N \times d_k}$, where d_k is the projected dimension.

With denotation of $\mathbf{q}_i, \mathbf{k}_j \in \mathbb{R}^{d_k}$ as the specific query and key of one (variante) token, we notice that each entry of the pre-Softmax scores is formulated as $\mathbf{A}_{i,j} = (\mathbf{Q}\mathbf{K}^\top / \sqrt{d_k})_{i,j} \propto \mathbf{q}_i^\top \mathbf{k}_j$. Since each token is previously normalized on its feature dimension, the entries can somewhat reveal the variate-wise correlation and the whole score map $\mathbf{A} \in \mathbb{R}^{N \times N}$ exhibits the multivariate correlations between paired variate tokens. Consequently, highly correlated variate will be more weighted for the next representation interaction with values \mathbf{V} . Based on this intuition, the proposed mechanism is believed to be more natural and interpretable for multivariate series forecasting. We further provide the visualization analysis of the score map in Section 4.3.

4 EXPERIMENTS

We thoroughly evaluate the proposed iTransformer on various time series forecasting applications, validate the generality of the proposed framework and further dive into the effects of inverting the duties of Transformer components for the specific time series dimensions.

Datasets We extensively include 6 real-world datasets in our experiments, including ETT, Weather, Electricity, Traffic used by Autoformer (Wu et al., 2021), Solar-Energy datasets proposed in LST-

Net (Lai et al., 2018), and PEMS evaluated in SCINet (Liu et al., 2022a). We also provide supplementary experiments of Market (6 subsets) in Appendix E.2, which records the minute-sampled server load of Alipay online transaction application with more than hundreds of variates. The detailed descriptions of datasets are provided in Appendix A.1.

4.1 FORECASTING RESULTS

In this section, we conduct extensive experiments to evaluate the forecasting performance of our proposed model together with advanced baselines.

Baselines We carefully choose 10 well-acknowledged forecasting models as our benchmark, including (1) Transformer-based methods: Informer (Li et al., 2021), Autoformer (Wu et al., 2021), FEDformer (Zhou et al., 2022), Stationary (Liu et al., 2022b), Crossformer (Zhang & Yan, 2023), PatchTST (Nie et al., 2023); (2) Linear-based methods: DLinear (Zeng et al., 2023), TiDE (Das et al., 2023); and (3) TCN-based methods: SCINet (Liu et al., 2022a), TimesNet (Wu et al., 2023).

Main results Comprehensive forecasting results are listed in Table 1 with the best in **red** and the second **underlined**. The lower MSE/MAE indicates the more accurate prediction result. The proposed iTransformer achieves consistent state-of-the-art performance. Notably, PatchTST as the previous best model on Electricity and Weather datasets, fails in many cases of PEMS, which can be attributed to the extremely fluctuating series of the dataset, and the patching mechanism of PatchTST may lose focus on specific locality to handle rapid fluctuation. By contrast, our proposed method aggregating the whole series variations for series representations can better cope with this situation. Besides, as the representative that explicitly captures multivariate correlations, the performance of Crossformer is still subpar to iTransformer, indicating the interaction of time-unaligned patches from different multivariate will bring about unnecessary noise for predictions. Therefore, native Transformer components are competent for temporal modeling and multivariate correlating, and the proposed inverted architecture can effectively tackle real-world time series forecasting scenarios.

Table 1: Multivariate forecasting results with prediction lengths $S \in \{12, 24, 36, 48\}$ for PEMS and $S \in \{96, 192, 336, 720\}$ for others. We fix the lookback length $T = 96$. All the results are averaged from all prediction lengths. Results of all prediction lengths are provided in Appendix E.2.

Models	iTransformer (Ours)	PatchTST (2023)	Crossformer (2023)	TiDE (2023)	TimesNet (2023)	DLinear	SCINet (2022a)	FEDformer (2022)	Stationary (2022b)	Autoformer (2021)	Informer (2021)
Metric	MSE MAE	MSE MAE	MSE MAE	MSE MAE	MSE MAE	MSE MAE	MSE MAE	MSE MAE	MSE MAE	MSE MAE	MSE MAE
ETT	0.383 0.407	<u>0.387</u> 0.407	0.942 0.684	0.611 0.550	0.414 0.427	0.559 0.515	0.954 0.723	0.437 0.449	0.526 0.516	0.450 0.459	4.431 1.729
Electricity	0.178 0.270	0.216 0.304	0.244 0.334	0.251 0.344	<u>0.192</u> 0.295	0.212 0.300	0.268 0.365	0.214 0.327	0.193 0.296	0.227 0.338	0.311 0.397
Traffic	0.428 0.282	0.555 0.362	<u>0.550</u> 0.304	0.760 0.473	0.620 0.336	0.625 0.383	0.804 0.509	0.610 0.376	0.624 0.340	0.628 0.379	0.764 0.416
Weather	0.258 0.279	<u>0.259</u> 0.281	0.259 0.315	0.271 0.320	0.259 0.287	0.265 0.317	0.292 0.363	0.309 0.360	0.288 0.314	0.338 0.382	0.634 0.548
Solar-Energy	0.233 0.262	0.270 0.307	0.641 0.639	0.347 0.417	0.301 0.319	0.330 0.401	0.282 0.375	0.291 0.381	0.261 0.381	0.885 0.711	<u>0.235</u> 0.280
PEMS	0.113 0.221	0.180 0.291	0.169 0.281	0.326 0.419	0.147 0.248	0.278 0.375	<u>0.114</u> 0.224	0.213 0.327	0.147 0.249	0.667 0.601	0.171 0.274

4.2 ITRANSFORMERS GENERALITY

In this section, we evaluate *iTransformers* by applying our framework to Transformer and its variants, which generally address the quadratic complexity of the self-attention mechanism, including Reformer (Kitaev et al., 2020), Informer (Li et al., 2021), Flowformer (Wu et al., 2022) and FlashAttention (Dao et al., 2022). Surprising and promising discoveries are exhibited, indicating the simple inverted perspective can enhance Transformer-based forecasters with promoted performance with efficiency, generalization on unseen variates, and better utilization of historical observations.

Performance promotion We evaluate Transformers and the corresponding iTransformers with the reported performance promotions in Table 2. It is notable that the framework consistently improves various Transformers. Overall, it achieves averaged **38.9%** promotion on Transformer, **36.1%** on Reformer, **28.5%** on Informer, **16.8%** on Flowformer and **32.2%** on Flashformer, revealing the previous improper usage of the Transformer architecture on time series forecasting. Moreover, since the attention mechanism is adopted on the variate dimension in our inverted structure, the introduction of efficient attentions with linear complexity essentially addresses the efficiency problem due to

numerous variates, which is prevalent in real-world applications but can be resource-consuming for Channel Independent (Nie et al., 2023). Therefore, the idea of iTransformer can be widely practiced on Transformer-based forecasters to take advantage of booming efficient attention mechanisms.

Table 2: Performance promotion obtained by our inverted framework. Flashformer means Transformer equipped with hardware-accelerated FlashAttention (Dao et al., 2022). We report the average performance and the relative MSE reduction (Promotion). Full results can be found in Appendix E.2.

Models		Transformer (2017)		Reformer (2020)		Informer (2021)		Flowformer (2022)		Flashformer (2022)	
Metric		MSE	MAE	MSE	MAE	MSE	MAE	MSE	MAE	MSE	MAE
Electricity	Original	0.277	0.372	0.338	0.422	0.311	0.397	0.267	0.359	0.285	0.377
	+Inverted	0.178	0.270	0.208	0.301	0.216	0.311	0.210	0.293	0.206	0.291
	Promotion	35.6%	27.4%	38.4%	28.7%	30.5%	21.6%	21.3%	18.6%	27.8%	22.9%
Traffic	Original	0.665	0.363	0.741	0.422	0.764	0.416	0.750	0.421	0.658	0.356
	+Inverted	0.428	0.282	0.647	0.370	0.662	0.380	0.524	0.355	0.492	0.333
	Promotion	35.6%	22.3%	12.7%	12.3%	13.3%	8.6%	30.1%	15.6%	25.2%	6.4%
Weather	Original	0.657	0.572	0.803	0.656	0.634	0.548	0.286	0.308	0.659	0.574
	+Inverted	0.258	0.279	0.248	0.292	0.271	0.330	0.266	0.285	0.262	0.282
	Promotion	60.2%	50.8%	69.2%	55.5%	57.3%	39.8%	7.2%	7.7%	60.2%	50.8%

Variate generalization By inverting vanilla Transformers, it is notable that the models are empowered with the generalization capability on unseen variates. Firstly, benefiting from the flexibility of the number of input tokens, the amount of variate channels is no longer restricted and thus feasible to vary from training and inference. Besides, feed-forward networks are identically applied on independent variate tokens in iTransformer. As aforementioned, the neurons as filters learn the intrinsic patterns of any time series, which are inclined to be shared among distinct variates.

To verify the hypothesis, we compare iTransformer with an alternative generalizing strategy: Channel Independent, which compulsively adopts one shared Transformer to learn the pattern of all variates. We train models with only 20% of variates from each dataset and directly test the trained model on all variates without fine-tuning. As shown in Figure 5, while the generalization error by Channel Independent (CI-Transformers) can increase tremendously, iTransformers present a much smaller increase in forecasting error. The difference lies in that the temporal interactions conducted in the attention of vanilla Transformer, well-acknowledged as a data-dependent module, can be less transferable for unseen series variations, for which feed-forward networks can be more proficient.

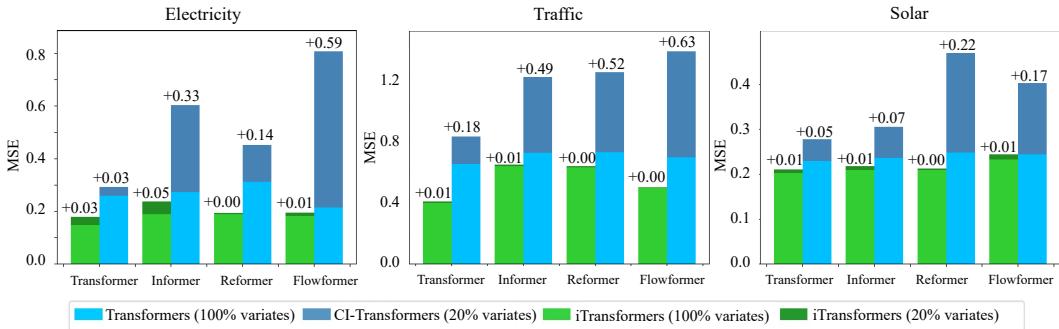


Figure 5: Performance of generalization on unseen variates. iTransformers can be naturally trained with 20% variates and accomplish forecast on all variates with hardly increased error, while Transformers with Channel Independence bring about significantly increased error.

Increasing lookback length Several previous works have observed the phenomenon that the forecasting performance does not necessarily improve with the increasing of lookback length on Transformers (Nie et al., 2023; Zeng et al., 2023), which can be attributed to the distracted attention on the growing input. However, the desired performance improvement is generally held on linear forecasts, theoretically supported by statistical methods (Box & Jenkins, 1968) with enlarged historical

information to be utilized. Since the duties of the attention and feed-forward network are inverted, we evaluate the performance of Transformers and iTransformer in Figure 6 with increased lookback length. The results surprisingly verify the rationality of leveraging MLPs on the temporal dimension such that Transformers can benefit from the extended lookback window for more precise predictions.

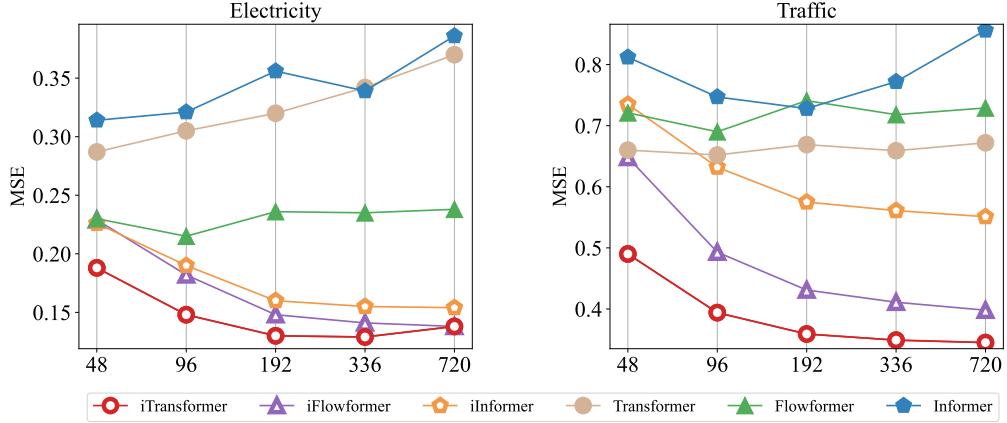


Figure 6: Forecasting performance with the lookback length $T \in \{48, 96, 192, 336, 720\}$ and fixed prediction length $S = 96$. While the performance of Transformer-based forecasters do not necessarily benefit from the increased lookback length, our inverted framework empowers the vanilla Transformer and its variants with improved performance on the enlarged lookback window.

4.3 MODEL ANALYSIS

Ablation study To verify the rational business of Transformer components, we provide detailed ablations covering both replacing components (Replace) and removing components (w/o) experiments. The results are listed in Table 3. iTransformer that utilizes attention on variate dimension and feed-forward on temporal dimension generally achieves the best performance. Notably, the performance of vanilla Transformer (the third row) performs the worst among these designs, indicating the disaccord of responsibility when the conventional architecture is adopted on time series forecasting.

Table 3: Ablations on iTransformer. We replace different components on the respective dimension to learn multivariate correlations (Variate) and series representations (Temporal), in addition to component removal. The average results of all predicted lengths are listed here.

Design	Variate	Temporal	Electricity		Traffic		Weather		Solar-Energy	
			MSE	MAE	MSE	MAE	MSE	MAE	MSE	MAE
iTransformer	Attention	FFN	0.178	0.270	0.428	0.282	0.258	0.278	0.233	0.262
Replace	Attention	Attention	0.193	0.293	0.913	0.500	0.255	0.280	0.261	0.291
	FFN	Attention	0.202	0.300	0.863	0.499	0.258	0.283	0.285	0.317
	FFN	FFN	0.182	0.287	0.599	0.348	0.248	0.274	0.269	0.287
w/o	Attention	w/o	0.189	0.278	0.456	0.306	0.261	0.281	0.258	0.289
	w/o	FFN	0.193	0.276	0.461	0.294	0.265	0.283	0.261	0.283

Analysis of series representations To further validate the claim that feed-forward networks are more favored to extract the series representations. We conduct representation analysis based on the centered kernel alignment (CKA) similarity (Kornblith et al., 2019). A higher CKA indicates more similar representations. For Transformer variants and iTransformers, we calculate the CKA between the output features of the first and the last block. Notably, previous works have demonstrated that time series forecasting, as a low-level generative task, prefers the higher CKA similarity (Wu et al., 2023; Dong et al., 2023) for the better performance. As shown in Figure 7, a clear line of division is exhibited, implying that iTransformers have learned more appropriate series representations by inverting the dimension and thus achieve more accurate predictions. The results also advocate inverting Transformer deserves a fundamental renovation of the forecasting backbone.

Analysis of multivariate correlations By assigning the duty of multivariate correlation to the attention mechanism, the learned map is able to enjoy enhanced interpretability. We present the case visualization on series from Soloar-Energy in Figure 7, which has distinct correlations in the lookback and future windows. It can be clearly observed that in the shallow attention layer, the learned map shares lots of similarities to the correlations of raw input series. As it dives into deeper layers, the learned map become gradually alike to the correlations of future series, which validates the inverted operation empowers interpretable attention for correlating, and the processes of encoding the past and decoding for the future are essentially conducted in series representations during feed-forwarding.

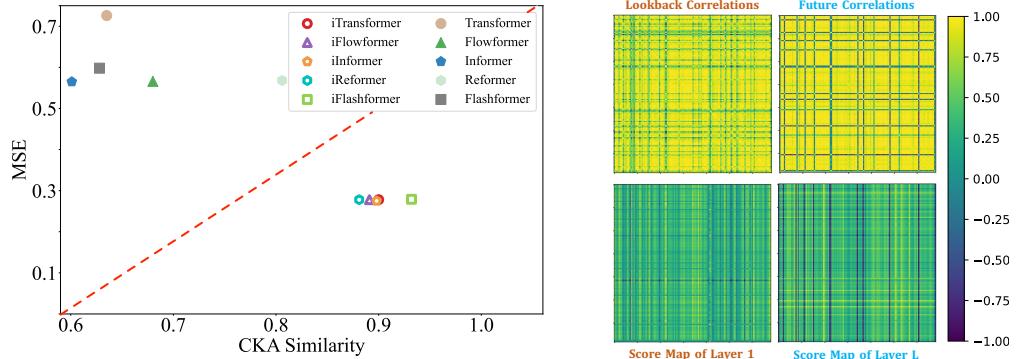


Figure 7: Analysis of series representations and multivariate correlations. Left: MSE and CKA similarity of representations comparison between Transformers and iTransformers. A higher CKA similarity indicates more favored representations for accurate predictions. Right: A case visualization of multivariate correlations of raw time series and the learned score maps by inverted self-attention.

Efficient training strategy In the proposed iTransformer, due to the quadratic complexity of self-attention, it can be overwhelming for training when variates grow numerous. In addition to efficient attention mechanisms, we propose a novel training strategy for high-dimensional multivariate series by taking advantage of previously demonstrated variate generation capability. Concretely, we randomly choose part of the variates in each batch and only train the model with selected variates. Since the number of variate channels is flexible because of our inverting, the model can predict all the variates for predictions. As shown in Figure 8, the performance of our proposed strategy is still comparable with full-variate training, while the memory footprint can be reduced significantly.

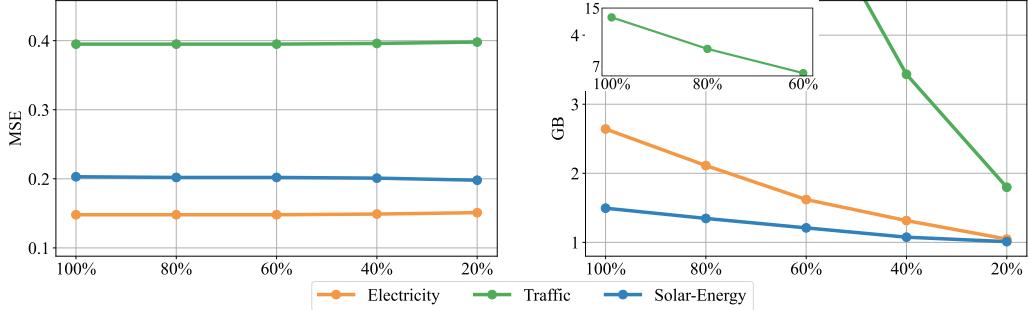


Figure 8: Analysis of the proposed training strategy. While the performance (Left) remains stable on partially trained variates of each batch with the sampled ratio in $\{20\%, 40\%, 60\%, 80\%, 100\%\}$, the memory footprint (Right) of the training process can be cut off significantly.

5 CONCLUSION AND FUTURE WORK

In this paper, we remark that the vanilla architecture of Transformer is not appropriate to discover essential series representations and multivariate correlations for time series forecasting. Therefore, we propose iTransformer as a fundamental backbone minimally adapted for time series, with native components naturally addressing the inverted dimensions. Experimentally, iTransformer achieves state-of-the-art performance and exhibits remarkable framework generality supported by promising analysis. In the future, we will explore iTransformers for extensive time series analysis tasks.

REFERENCES

- Jimmy Lei Ba, Jamie Ryan Kiros, and Geoffrey E. Hinton. Layer normalization. <https://arxiv.org/pdf/1607.06450.pdf>, 2016.
- Shaojie Bai, J Zico Kolter, and Vladlen Koltun. An empirical evaluation of generic convolutional and recurrent networks for sequence modeling. *arXiv preprint arXiv:1803.01271*, 2, 2018.
- George EP Box and Gwilym M Jenkins. Some recent advances in forecasting and control. *Journal of the Royal Statistical Society. Series C (Applied Statistics)*, 17(2):91–109, 1968.
- Tom Brown, Benjamin Mann, Nick Ryder, Melanie Subbiah, Jared D Kaplan, Prafulla Dhariwal, Arvind Neelakantan, Pranav Shyam, Girish Sastry, Amanda Askell, et al. Language models are few-shot learners. *NeurIPS*, 2020.
- Tri Dao, Dan Fu, Stefano Ermon, Atri Rudra, and Christopher Ré. Flashattention: Fast and memory-efficient exact attention with io-awareness. *NeurIPS*, 2022.
- Abhimanyu Das, Weihao Kong, Andrew Leach, Rajat Sen, and Rose Yu. **Long-term forecasting with tide: Time-series dense encoder**. *arXiv preprint arXiv:2304.08424*, 2023.
- Jiaxiang Dong, Haixu Wu, Haoran Zhang, Li Zhang, Jianmin Wang, and Mingsheng Long. Simmtm: A simple pre-training framework for masked time-series modeling. *arXiv preprint arXiv:2302.00861*, 2023.
- Alexey Dosovitskiy, Lucas Beyer, Alexander Kolesnikov, Dirk Weissenborn, Xiaohua Zhai, Thomas Unterthiner, Mostafa Dehghani, Matthias Minderer, Georg Heigold, Sylvain Gelly, et al. An image is worth 16x16 words: Transformers for image recognition at scale. *ICLR*, 2021.
- Vijay Ekambaram, Arindam Jati, Nam Nguyen, Phanwadee Sinthong, and Jayant Kalagnanam. Tsmixer: Lightweight mlp-mixer model for multivariate time series forecasting. *KDD*, 2023.
- Kurt Hornik. Approximation capabilities of multilayer feedforward networks. *Neural networks*, 4(2): 251–257, 1991.
- Jared Kaplan, Sam McCandlish, Tom Henighan, Tom B Brown, Benjamin Chess, Rewon Child, Scott Gray, Alec Radford, Jeffrey Wu, and Dario Amodei. Scaling laws for neural language models. *arXiv preprint arXiv:2001.08361*, 2020.
- Taesung Kim, Jinhee Kim, Yunwon Tae, Cheonbok Park, Jang-Ho Choi, and Jaegul Choo. Reversible instance normalization for accurate time-series forecasting against distribution shift. *ICLR*, 2021.
- Diederik P. Kingma and Jimmy Ba. Adam: A method for stochastic optimization. *ICLR*, 2015.
- Nikita Kitaev, Łukasz Kaiser, and Anselm Levskaya. Reformer: The efficient transformer. *ICLR*, 2020.
- Simon Kornblith, Mohammad Norouzi, Honglak Lee, and Geoffrey E. Hinton. Similarity of neural network representations revisited. *ICML*, 2019.
- Guokun Lai, Wei-Cheng Chang, Yiming Yang, and Hanxiao Liu. Modeling long-and short-term temporal patterns with deep neural networks. *SIGIR*, 2018.
- Jianxin Li, Xiong Hui, and Wancai Zhang. Informer: Beyond efficient transformer for long sequence time-series forecasting. *arXiv: 2012.07436*, 2021.
- Zhe Li, Shiyi Qi, Yiduo Li, and Zenglin Xu. Revisiting long-term time series forecasting: An investigation on linear mapping. *arXiv preprint arXiv:2305.10721*, 2023.
- Minhao Liu, Ailing Zeng, Muxi Chen, Zhijian Xu, Qiuxia Lai, Lingna Ma, and Qiang Xu. Scinet: time series modeling and forecasting with sample convolution and interaction. *NeurIPS*, 2022a.
- Yong Liu, Haixu Wu, Jianmin Wang, and Mingsheng Long. Non-stationary transformers: Rethinking the stationarity in time series forecasting. *NeurIPS*, 2022b.

-
- Yuqi Nie, Nam H Nguyen, Phanwadee Sinthong, and Jayant Kalagnanam. A time series is worth 64 words: Long-term forecasting with transformers. *ICLR*, 2023.
- Boris N Oreshkin, Dmitri Carpow, Nicolas Chapados, and Yoshua Bengio. N-BEATS: Neural basis expansion analysis for interpretable time series forecasting. *ICLR*, 2019.
- Adam Paszke, S. Gross, Francisco Massa, A. Lerer, James Bradbury, Gregory Chanan, Trevor Killeen, Z. Lin, N. Gimelshein, L. Antiga, Alban Desmaison, Andreas Köpf, Edward Yang, Zach DeVito, Martin Raison, Alykhan Tejani, Sasank Chilamkurthy, Benoit Steiner, Lu Fang, Junjie Bai, and Soumith Chintala. Pytorch: An imperative style, high-performance deep learning library. *NeurIPS*, 2019.
- Syama Sundar Rangapuram, Matthias W Seeger, Jan Gasthaus, Lorenzo Stella, Yuyang Wang, and Tim Januschowski. Deep state space models for time series forecasting. *NeurIPS*, 2018.
- David Salinas, Valentin Flunkert, Jan Gasthaus, and Tim Januschowski. Deepar: Probabilistic forecasting with autoregressive recurrent networks. *International Journal of Forecasting*, 36(3): 1181–1191, 2020.
- Ilya O Tolstikhin, Neil Houlsby, Alexander Kolesnikov, Lucas Beyer, Xiaohua Zhai, Thomas Unterthiner, Jessica Yung, Andreas Steiner, Daniel Keysers, Jakob Uszkoreit, et al. Mlp-mixer: An all-mlp architecture for vision. *NeurIPS*, 2021.
- Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N Gomez, Łukasz Kaiser, and Illia Polosukhin. Attention is all you need. *NeurIPS*, 2017.
- Haixu Wu, Jiehui Xu, Jianmin Wang, and Mingsheng Long. Autoformer: Decomposition transformers with Auto-Correlation for long-term series forecasting. *NeurIPS*, 2021.
- Haixu Wu, Jialong Wu, Jiehui Xu, Jianmin Wang, and Mingsheng Long. Flowformer: Linearizing transformers with conservation flows. *ICML*, 2022.
- Haixu Wu, Tengge Hu, Yong Liu, Hang Zhou, Jianmin Wang, and Mingsheng Long. Timesnet: Temporal 2d-variation modeling for general time series analysis. *ICLR*, 2023.
- Ailing Zeng, Muxi Chen, Lei Zhang, and Qiang Xu. Are transformers effective for time series forecasting? *AAAI*, 2023.
- Yunhao Zhang and Junchi Yan. Crossformer: Transformer utilizing cross-dimension dependency for multivariate time series forecasting. *ICLR*, 2023.
- Zheng Zhao, Weihai Chen, Xingming Wu, Peter CY Chen, and Jingmeng Liu. Lstm network: a deep learning approach for short-term traffic forecast. *IET Intelligent Transport Systems*, 11(2):68–75, 2017.
- Tian Zhou, Ziqing Ma, Qingsong Wen, Xue Wang, Liang Sun, and Rong Jin. FEDformer: Frequency enhanced decomposed transformer for long-term series forecasting. *ICML*, 2022.

A IMPLEMENTATION DETAILS

A.1 DATASETS DESCRIPTIONS

We conduct experiments on 12 real-world datasets to evaluate the performance of the proposed iTransformer including: (1) Electricity (Wu et al., 2021) records the hourly electricity consumption data of 321 clients. (2) ETT (Li et al., 2021) contains 7 factors of electricity transformer from July 2016 to July 2018, which is recorded by hour. (3) Traffic (Wu et al., 2021) collects hourly road occupancy rates measured by 862 sensors of San Francisco Bay area freeways from January 2015 to December 2016. (4) Solar-Energy (Lai et al., 2018) records the solar power production of 137 PV plants in 2006, which is sampled every 10 minutes. (5) Weather (Wu et al., 2021) includes 21 meteorological factors collected every 10 minutes from the Weather Station of the Max Planck Biogeochemistry Institute in 2020. (6) PEMS (Liu et al., 2022a) contains public traffic network data in California collected by 5-minute windows with 358 variates.

Apart from the public datasets widely used as forecasting benchmarks, we also collect a set of Market datasets of a real-world application, which records the minute-sampled server load of Alipay online transactions between January 30th, 2023, and April 9th, 2023 with the number of variates varied from 285 to 759. It includes 6 sub-datasets, which are divided according to diverse transaction domains.

We follow the same data processing protocol and forecasting length settings used in TimesNet (Wu et al., 2023) and SCINet (Liu et al., 2022a). For ETT, Weather, Electricity, Solar-Energy, and Traffic datasets, we fix the length of the lookback series as 96, and the prediction length varies in $\{96, 192, 336, 720\}$. For PEMS dataset, we fix the lookback length as 12, and the prediction length varies in $\{12, 24, 36, 48\}$. For the Market dataset, the lookback window contains the past one day observations with 144 time points and the forecasting length varies in $\{12, 24, 72, 144\}$. The details of all the datasets are provided in Table 4.

Table 4: Detailed dataset descriptions. *Dim* denotes the variate number of each dataset. *Dataset Size* denotes the total number of time points in (Train, Validation, Test) split respectively. *Prediction Length* denotes the future time points to be predict and four prediction settings are included in each dataset. *Frequency* denotes the sampling interval of time points.

Dataset	Dim	Prediction Length	Dataset Size	Frequency	Information
ETT	7	{96, 192, 336, 720}	(8545, 2881, 2881)	15min	Electricity
Weather	21	{96, 192, 336, 720}	(36792, 5271, 10540)	10min	Weather
Solar-Energy	137	{96, 192, 336, 720}	(36601, 5161, 10417)	10min	Energy
Electricity	321	{96, 192, 336, 720}	(18317, 2633, 5261)	Hourly	Electricity
Traffic	862	{96, 192, 336, 720}	(12185, 1757, 3509)	Hourly	Transportation
PEMS	358	{12, 24, 48, 96}	(15701, 5216, 434)	5min	Transportation
Market-Merchant	285	{12, 24, 72, 144}	(7045, 1429, 1429)	10min	Transaction
Market-Wealth	485	{12, 24, 72, 144}	(7045, 1429, 1429)	10min	Transaction
Market-Finance	405	{12, 24, 72, 144}	(7045, 1429, 1429)	10min	Transaction
Market-Terminal	307	{12, 24, 72, 144}	(7045, 1429, 1429)	10min	Transaction
Market-Payment	759	{12, 24, 72, 144}	(7045, 1429, 1429)	10min	Transaction
Market-Customer	395	{12, 24, 72, 144}	(7045, 1429, 1429)	10min	Transaction

A.2 IMPLIEMNTATION DETAILS

All the experiments are implemented in PyTorch (Paszke et al., 2019) and conducted on a single NVIDIA P100 16GB GPU. We utilize ADAM (Kingma & Ba, 2015) with an initial learning rate in $\{10^{-3}, 5 \times 10^{-4}, 10^{-4}\}$ and L2 loss for the model optimization. The batch size is uniformly set to 32 and the number of training epochs is fixed to 10. We set the number of inverted Transformer blocks in our proposed model $L \in \{2, 3, 4\}$. The dimension of series representations D is set from $\{256, 512\}$. All the compared baseline models that we reproduced are implemented based on the benchmark of

TimesNet (Wu et al., 2023) Repository, which is fairly built on the configurations provided by each model’s original paper or official code. Since several baselines have adopted Series Stationarization of Non-stationary Transformers (Liu et al., 2022b) while others do not, we equip all models with the method for a fair comparison. And we also provide the pseudo-code of iTransformer in Algorithm 1.

Algorithm 1 iTransformer - Overall Architecture.

Require: Input lookback time series $\mathbf{X} \in \mathbb{R}^{T \times N}$; input Length T ; predicted length S ; variables number N ; token dimension D ; iTransformer block number L .

- 1: $\mathbf{X} = \mathbf{X}.\text{transpose}$ $\triangleright \mathbf{X} \in \mathbb{R}^{N \times T}$
- 2: \triangleright Multi-layer Perceptron works on the last dimension to embed series into variate tokens.
- 3: $\mathbf{H}^0 = \text{MLP}(\mathbf{X})$ $\triangleright \mathbf{H}^0 \in \mathbb{R}^{N \times D}$
- 4: **for** l **in** $\{1, \dots, L\}$: \triangleright Run through iTransformer blocks.
- 5: \triangleright Self-attention layer is applied on variate tokens.
- 6: $\mathbf{H}^{l-1} = \text{LayerNorm}(\mathbf{H}^{l-1} + \text{Self-Attn}(\mathbf{H}^{l-1}))$ $\triangleright \mathbf{H}^{l-1} \in \mathbb{R}^{N \times D}$
- 7: \triangleright Feed-forward network is utilized for series representations, broadcasting to each token.
- 8: $\mathbf{H}^l = \text{LayerNorm}(\mathbf{H}^{l-1} + \text{Feed-Forward}(\mathbf{H}^{l-1}))$ $\triangleright \mathbf{H}^l \in \mathbb{R}^{N \times D}$
- 9: \triangleright LayerNorm is adopted on series representations to reduce variates discrepancies.
- 10: **End for**
- 11: $\hat{\mathbf{Y}} = \text{MLP}(\mathbf{H}^L)$ \triangleright Project tokens back to predicted series, $\hat{\mathbf{Y}} \in \mathbb{R}^{N \times S}$
- 12: $\hat{\mathbf{Y}} = \hat{\mathbf{Y}}.\text{transpose}$ $\triangleright \hat{\mathbf{Y}} \in \mathbb{R}^{S \times N}$
- 13: **Return** $\hat{\mathbf{Y}}$ \triangleright Return the prediction result $\hat{\mathbf{Y}}$

B ABLATION STUDIES

To elaborate on the rational business of Transformer components, we conduct detailed ablations covering replacing components (Replace) and removing components (w/o). Since the average results are listed in Table 3 due to the paper limit, we provide the detailed results and analysis here.

As shown in Table 5, among various architectural designs, iTransformer learns multivariate correlations by self-attention and encodes series representations by the feed-forward network and generally exhibit superior performance. Nevertheless, the vanilla Transformer architecture that leverages feed-forward on variate representations and self-attention on time points can lead to the worst performance, indicating the misuse of Transformer components on time series forecasting. It is also notable that applying pure feed-forward networks on both variate and temporal dimensions can also lead to good performance in Weather forecasting. However, we find the phenomenon only happens in the dataset with a relatively small number of time series (21 variates). With the increasing number of variates and difficulty in predicting the time series, the importance of capturing multivariate correlations is ever more highlighted. Our proposed iTransformer employing the self-attention module to disentangle the correlations between variate tokens proves to be more powerful than aggregating the variate representations by feed-forward networks, thereby further boosting the performance on challenging datasets with more variates and enhancing the interpretability.

C HYPERPARAMETER SENSITIVITY

We evaluate the hyperparameter sensitivity of iTransformer with respect to the following factors: the learning rate lr , the number of Transformer block L , and the hidden dimension D of variate tokens. The results are shown in Figure 9. We find that the learning rate, as the most common influencing factor, should be carefully selected when the number of variates is large (ECL, Traffic). And the block number and hidden dimension are not essentially favored to be as large as possible in iTransformer.

Table 5: Full results of the ablation on iTransformer. We apply different components on the respective dimension to learn multivariate correlations (Variate) and series representations (Temporal), in addition to removing the specific component of Transformer.

Design	Variate	Temporal	Prediction	Electricity		Traffic		Weather		Solar-Energy	
				Lengths	MSE	MAE	MSE	MAE	MSE	MAE	MSE
iTransformer	Attention	FFN	96	0.148	0.240	0.395	0.268	0.174	0.214	0.203	0.237
			192	0.162	0.253	0.417	0.276	0.221	0.254	0.233	0.261
			336	0.178	0.269	0.433	0.283	0.278	0.296	0.248	0.273
			720	0.225	0.317	0.467	0.302	0.358	0.349	0.249	0.275
			Avg	0.178	0.270	0.428	0.282	0.258	0.279	0.233	0.262
Replace	Attention	Attention	96	0.161	0.263	1.021	0.581	0.168	0.213	0.227	0.270
			192	0.180	0.280	0.834	0.447	0.217	0.256	0.255	0.292
			336	0.194	0.296	0.906	0.493	0.277	0.299	0.279	0.301
			720	0.238	0.331	0.892	0.477	0.356	0.351	0.283	0.300
			Avg	0.193	0.293	0.913	0.500	0.255	0.280	0.261	0.291
	FFN	Attention	96	0.169	0.270	0.907	0.540	0.176	0.221	0.247	0.299
			192	0.189	0.292	0.839	0.489	0.224	0.261	0.275	0.305
			336	0.204	0.304	0.248	0.364	0.279	0.301	0.317	0.337
			720	0.245	0.335	1.059	0.606	0.354	0.347	0.301	0.329
			Avg	0.202	0.300	0.863	0.499	0.258	0.283	0.285	0.317
	FFN	FFN	96	0.159	0.261	0.606	0.342	0.162	0.207	0.237	0.277
			192	0.171	0.271	0.559	0.342	0.211	0.252	0.273	0.293
			336	0.187	0.287	0.569	0.348	0.270	0.293	0.284	0.287
			720	0.211	0.307	0.664	0.359	0.349	0.345	0.284	0.289
			Avg	0.182	0.287	0.599	0.348	0.248	0.274	0.269	0.287
w/o	Attention	w/o	96	0.163	0.254	0.427	0.296	0.177	0.219	0.226	0.266
			192	0.174	0.263	0.446	0.300	0.226	0.259	0.255	0.288
			336	0.191	0.280	0.459	0.306	0.281	0.298	0.275	0.301
			720	0.228	0.315	0.492	0.324	0.359	0.249	0.275	0.301
			Avg	0.189	0.278	0.456	0.306	0.261	0.281	0.258	0.289
	w/o	FFN	96	0.169	0.253	0.437	0.283	0.183	0.220	0.228	0.263
			192	0.177	0.261	0.449	0.287	0.231	0.262	0.261	0.283
			336	0.194	0.278	0.464	0.294	0.285	0.300	0.279	0.294
			720	0.233	0.311	0.496	0.313	0.362	0.350	0.276	0.291
			Avg	0.193	0.276	0.461	0.294	0.265	0.283	0.261	0.283

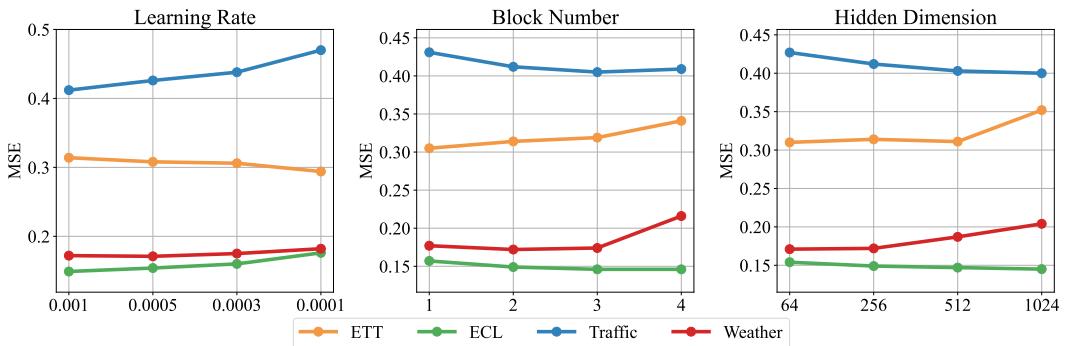


Figure 9: Hyperparameter sensitivity with respect to the learning rate, the number of Transformer block, and the hidden dimension of variate tokens. The results are recorded with the lookback window length $T = 96$ and the forecast window length $S = 96$.

D SHOWCASES

D.1 VISUALIZATION OF MULTIVARIATE CORRELATIONS

By using the attention mechanism on multiple variate tokens, the resulting learned map can become more interpretable. To present an intuitive understanding of the multivariate correlations. In Figure 10, we provide three randomly chosen case visualizations of the time series from Solar-Energy. Each case is divided by the dotted line. On the top part of each case, we provide the variate correlations calculated on the raw series by Pearson Correlation coefficient as the following equation:

$$\rho_{xy} = \frac{\sum_i(x_i - \bar{x})(y_i - \bar{y})}{\sqrt{\sum_i(x_i - \bar{x})^2}\sqrt{\sum_i(y_i - \bar{y})^2}},$$

where $x_i, y_i \in \mathbb{R}$ run through all time points of the paired variates to be correlated. All the cases have distinct multivariate correlations in the lookback and future time series because the dataset exhibits obvious seasonal change in the daytime and night. On the bottom part of each case, we provide the learned pre-Softmax maps of the self-attention module in both the first and the last but one layers. As we observe in the shallow attention layer, we can find that the learned map is similar to the correlations of the raw lookback series. As we go deeper into the layers, the learned map gradually becomes more similar to the correlations of the future series to be predicted. This demonstrates that the inverted operation allows for interpretable attention in correlating, and that encoding of the past and decoding for the future are conducted through series representations during layer stacking.

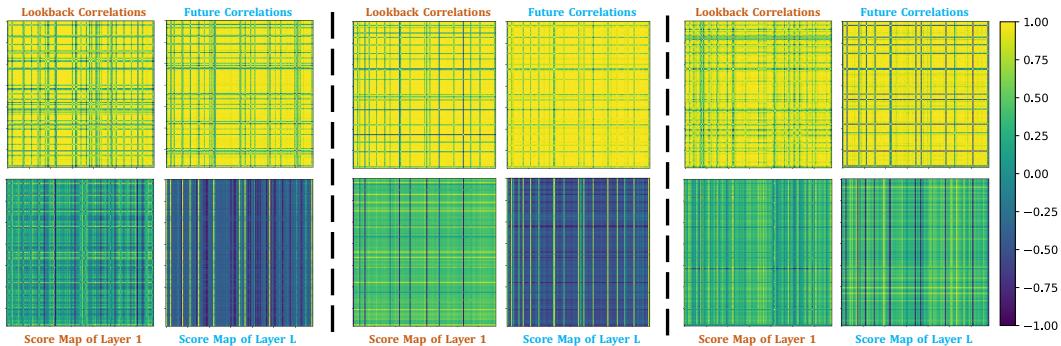


Figure 10: Multivariate correlations of the lookback series and future series and the learned score maps by inverted self-attention of different layers. Cases all come from the Solar-Energy dataset.

D.2 VISUALIZATION OF PREDICTION RESULTS

To provide a clear comparison among different models, we list supplementary prediction showcases of three representative datasets in Figures 12, 11, and 13, which are given by the following models: iTransfomer, PatchTST (Nie et al., 2023), DLinear (Zeng et al., 2023), Crossformer (Zhang & Yan, 2023), Autoformer (Wu et al., 2021), Transformer (Vaswani et al., 2017). Among the various models, iTransformer predicts the most precise future series variations and exhibits superior performance.

E FULL RESULTS

E.1 FULL FRAMEWORK PROMOTION RESULTS

We apply the proposed iTransformers framework to five Transformer and its variants: Transformer (Vaswani et al., 2017), Reformer (Kitaev et al., 2020), Informer (Li et al., 2021), Flowformer (Wu et al., 2022), Flashformer (Dao et al., 2022). The averaged results are shown in Table 2 due to the limited pages. We provide the supplementary forecasting results in Table 6. The results demonstrate that our iTransformers framework can consistently promote these Transformer variants, and take advantage of the booming efficient attention mechanisms.

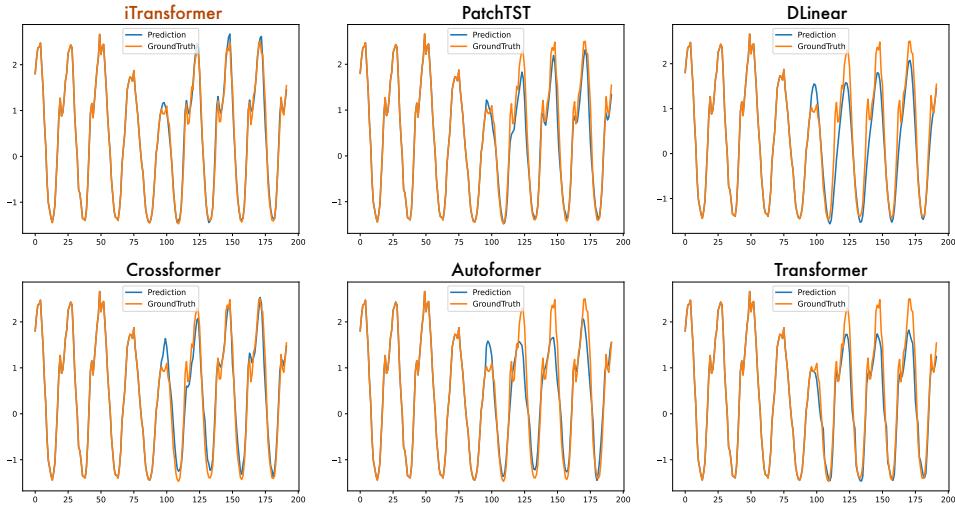


Figure 11: Visualization of input-96-predict-96 results on the Traffic dataset.

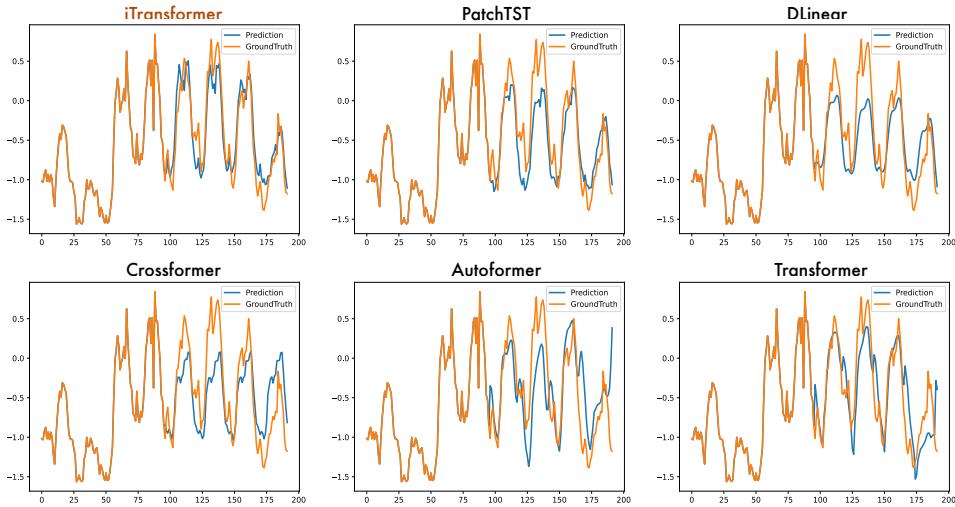


Figure 12: Visualization of input-96-predict-96 results on the Electricity dataset.

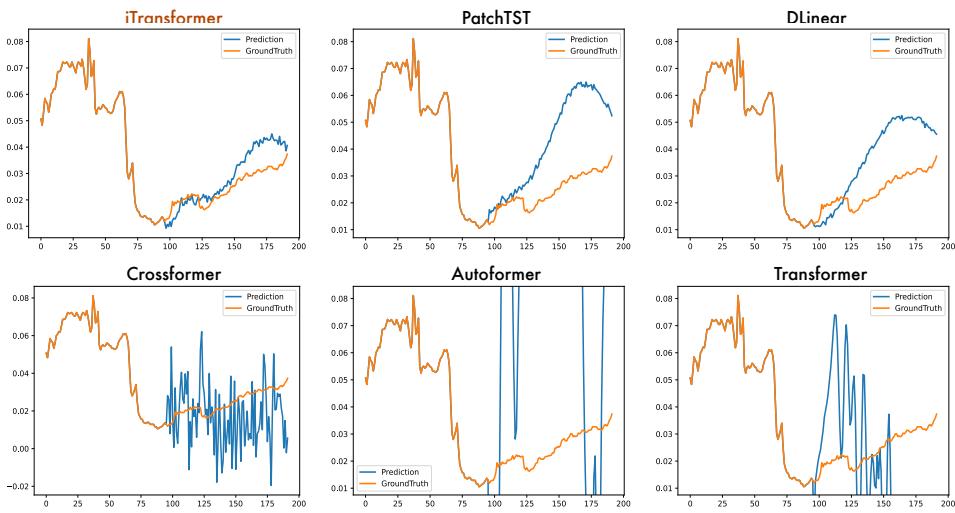


Figure 13: Visualization of input-96-predict-96 results on the Weather dataset.

Table 6: Full results of Transformers with our inverted framework. Flashformer means Transformer equipped with hardware-accelerated FlashAttention (Dao et al., 2022).

Models		Transformer (2017)		Reformer (2020)		Informer (2021)		Flowformer (2022)		Flashformer (2022)		
Metric		MSE	MAE	MSE	MAE	MSE	MAE	MSE	MAE	MSE	MAE	
Electricity	Original	96	0.260	0.358	0.312	0.402	0.274	0.368	0.215	0.320	0.259	0.357
		192	0.266	0.367	0.348	0.433	0.296	0.386	0.259	0.355	0.274	0.374
		336	0.280	0.375	0.350	0.433	0.300	0.394	0.296	0.383	0.310	0.396
		720	0.302	0.386	0.340	0.420	0.373	0.439	0.296	0.380	0.298	0.383
	Avg	0.277	0.372	0.338	0.422	0.311	0.397	0.267	0.359	0.285	0.377	
	+Inverted	96	0.148	0.240	0.182	0.275	0.190	0.286	0.183	0.267	0.178	0.265
		192	0.162	0.253	0.192	0.286	0.201	0.297	0.192	0.277	0.189	0.276
		336	0.178	0.269	0.210	0.304	0.218	0.315	0.210	0.295	0.207	0.294
		720	0.225	0.317	0.249	0.339	0.255	0.347	0.255	0.332	0.251	0.329
	Avg	0.178	0.270	0.208	0.301	0.216	0.311	0.210	0.293	0.206	0.291	
Traffic	Original	96	0.647	0.357	0.732	0.423	0.719	0.391	0.691	0.393	0.641	0.348
		192	0.649	0.356	0.733	0.420	0.696	0.379	0.729	0.419	0.648	0.358
		336	0.667	0.364	0.742	0.420	0.777	0.420	0.756	0.423	0.670	0.364
		720	0.697	0.376	0.755	0.432	0.864	0.472	0.825	0.449	0.673	0.354
	Avg	0.665	0.363	0.741	0.422	0.764	0.416	0.750	0.421	0.658	0.356	
	+Inverted	96	0.395	0.268	0.617	0.356	0.632	0.367	0.493	0.339	0.464	0.320
		192	0.417	0.276	0.629	0.361	0.641	0.370	0.506	0.345	0.479	0.326
		336	0.433	0.283	0.648	0.370	0.663	0.379	0.526	0.355	0.501	0.337
		720	0.467	0.302	0.694	0.394	0.713	0.405	0.572	0.381	0.524	0.350
	Avg	0.428	0.282	0.647	0.370	0.662	0.380	0.524	0.355	0.492	0.333	
Weather	Original	96	0.395	0.427	0.689	0.596	0.300	0.384	0.182	0.233	0.388	0.425
		192	0.619	0.560	0.752	0.638	0.598	0.544	0.250	0.288	0.619	0.560
		336	0.689	0.594	0.639	0.596	0.578	0.523	0.309	0.329	0.698	0.600
		720	0.926	0.710	1.130	0.792	1.059	0.741	0.404	0.385	0.930	0.711
	Avg	0.657	0.572	0.803	0.656	0.634	0.548	0.286	0.308	0.659	0.574	
	+Inverted	96	0.174	0.214	0.169	0.225	0.180	0.251	0.183	0.223	0.177	0.218
		192	0.221	0.254	0.213	0.265	0.244	0.318	0.231	0.262	0.229	0.261
		336	0.278	0.296	0.268	0.317	0.282	0.343	0.286	0.301	0.283	0.300
		720	0.358	0.349	0.340	0.361	0.377	0.409	0.363	0.352	0.359	0.251
	Avg	0.258	0.279	0.248	0.292	0.271	0.330	0.266	0.285	0.262	0.282	

E.2 FULL FORECASTING RESULTS

The full multivariate forecasting results are provided in the following section due to the space limitation of the main text. Table 7 contains the detailed results of all prediction lengths of the six well-acknowledged benchmarks. And Table 8 records the Market results for server load forecasting. Our proposed model achieves consistent state-of-the-art in real-world forecasting applications.

Table 7: Full results for the long-term forecasting task. We compare extensive competitive models under different prediction lengths following the setting of TimesNet (2023). The input sequence length is set to 96 for all baselines. Avg means the average results from all four prediction lengths.

Models	iTransformer (Ours)	PatchTST (2023)	Crossformer (2023)	SCINet (2022a)	TiDE (2023)	TimesNet (2023)	DLinear (2022)	FEDformer (2022b)	Stationary (2021)	Autoformer (2021)	Informer (2021)	
Metric	MSE	MAE	MSE MAE	MSE MAE	MSE MAE	MSE MAE	MSE MAE	MSE MAE	MSE MAE	MSE MAE	MSE MAE	
ETT	96	0.297	0.349	0.302 0.348	0.745 0.584	0.707 0.621	0.400 0.440	0.340 0.374	0.333 0.387	0.358 0.397	0.476 0.458	0.346 0.388
	192	0.380	0.400	0.388 0.400	0.877 0.656	0.860 0.689	0.528 0.509	0.402 0.414	0.477 0.476	0.429 0.439	0.512 0.493	0.456 0.452
	336	0.428	0.432	0.426 0.433	1.043 0.731	1.000 0.744	0.643 0.571	0.452 0.452	0.594 0.541	0.496 0.487	0.552 0.551	0.482 0.486
	720	0.427	0.445	0.431 0.446	1.104 0.763	1.249 0.838	0.874 0.679	0.462 0.468	0.831 0.657	0.463 0.474	0.562 0.560	0.515 0.511
	Avg	0.383	0.407	0.387 0.407	0.942 0.684	0.954 0.723	0.611 0.550	0.414 0.427	0.559 0.515	0.437 0.449	0.526 0.516	0.450 0.459
Electricity	96	0.148	0.240	0.195 0.285	0.219 0.314	0.247 0.345	0.237 0.329	0.168 0.272	0.197 0.282	0.193 0.308	0.169 0.273	0.201 0.317
	192	0.162	0.253	0.199 0.289	0.231 0.322	0.257 0.355	0.236 0.330	0.184 0.289	0.196 0.285	0.201 0.315	0.182 0.286	0.222 0.334
	336	0.178	0.269	0.215 0.305	0.246 0.337	0.269 0.369	0.249 0.344	0.198 0.300	0.209 0.301	0.214 0.329	0.200 0.304	0.231 0.338
	720	0.225	0.317	0.256 0.337	0.280 0.363	0.299 0.390	0.284 0.373	0.220 0.320	0.245 0.333	0.246 0.355	0.222 0.321	0.254 0.361
	Avg	0.178	0.270	0.216 0.304	0.244 0.334	0.268 0.365	0.251 0.344	0.192 0.295	0.212 0.300	0.214 0.327	0.193 0.296	0.227 0.338
Traffic	96	0.395	0.268	0.544 0.359	0.522 0.290	0.788 0.499	0.805 0.493	0.593 0.321	0.650 0.396	0.587 0.366	0.612 0.338	0.613 0.388
	192	0.417	0.276	0.540 0.354	0.530 0.293	0.789 0.505	0.756 0.474	0.617 0.336	0.598 0.370	0.604 0.373	0.613 0.340	0.616 0.382
	336	0.433	0.283	0.551 0.358	0.558 0.305	0.797 0.508	0.762 0.477	0.629 0.336	0.605 0.373	0.621 0.383	0.618 0.328	0.622 0.337
	720	0.467	0.302	0.586 0.375	0.589 0.328	0.841 0.523	0.719 0.449	0.640 0.350	0.645 0.394	0.626 0.382	0.653 0.355	0.660 0.408
	Avg	0.428	0.282	0.555 0.362	0.550 0.304	0.804 0.509	0.760 0.473	0.620 0.336	0.625 0.383	0.610 0.376	0.624 0.340	0.628 0.379
Weather	96	0.174	0.214	0.177 0.218	0.158 0.230	0.221 0.306	0.202 0.261	0.172 0.220	0.196 0.255	0.217 0.296	0.173 0.223	0.266 0.336
	192	0.221	0.254	0.225 0.259	0.206 0.277	0.261 0.340	0.242 0.298	0.219 0.261	0.237 0.296	0.276 0.336	0.245 0.285	0.307 0.367
	336	0.278	0.296	0.278 0.297	0.272 0.335	0.309 0.378	0.287 0.353	0.280 0.306	0.283 0.335	0.339 0.380	0.321 0.338	0.359 0.395
	720	0.358	0.349	0.354 0.348	0.398 0.418	0.377 0.427	0.351 0.386	0.365 0.359	0.345 0.381	0.403 0.428	0.414 0.410	0.419 0.428
	Avg	0.258	0.279	0.259 0.281	0.259 0.315	0.292 0.363	0.271 0.320	0.259 0.287	0.265 0.317	0.309 0.360	0.288 0.314	0.338 0.382
Solar-Energy	96	0.203	0.237	0.234 0.286	0.310 0.331	0.237 0.344	0.312 0.399	0.250 0.292	0.290 0.378	0.242 0.342	0.215 0.249	0.884 0.711
	192	0.233	0.261	0.267 0.310	0.734 0.725	0.280 0.380	0.339 0.416	0.296 0.318	0.320 0.398	0.285 0.380	0.254 0.272	0.834 0.692
	336	0.248	0.273	0.290 0.315	0.750 0.735	0.304 0.389	0.368 0.430	0.319 0.330	0.353 0.415	0.282 0.376	0.290 0.296	0.941 0.723
	720	0.249	0.275	0.289 0.317	0.769 0.765	0.308 0.388	0.370 0.425	0.338 0.337	0.356 0.413	0.357 0.427	0.285 0.295	0.882 0.717
	Avg	0.233	0.262	0.270 0.307	0.641 0.639	0.282 0.375	0.347 0.417	0.301 0.319	0.330 0.401	0.291 0.381	0.261 0.381	0.885 0.711
PEMS	12	0.071	0.174	0.099 0.216	0.090 0.203	0.066 0.172	0.178 0.305	0.085 0.192	0.122 0.243	0.126 0.251	0.081 0.188	0.272 0.385
	24	0.093	0.201	0.142 0.259	0.121 0.240	0.085 0.198	0.257 0.371	0.118 0.223	0.201 0.317	0.149 0.275	0.105 0.214	0.334 0.440
	48	0.125	0.236	0.211 0.319	0.202 0.317	0.127 0.238	0.379 0.463	0.155 0.260	0.333 0.425	0.227 0.348	0.154 0.257	1.032 0.782
	96	0.160	0.270	0.269 0.370	0.262 0.367	0.178 0.287	0.490 0.539	0.228 0.317	0.457 0.515	0.348 0.434	0.247 0.336	1.031 0.796
	Avg	0.113	0.221	0.180 0.291	0.169 0.281	0.114 0.224	0.326 0.419	0.147 0.248	0.278 0.375	0.213 0.327	0.147 0.249	0.667 0.601

Table 8: Full results of the Market dataset. We compare extensive competitive models on the real-world transaction forecasting task. Avg means the average results from all prediction lengths.

Models	iTransformer (Ours)	PatchTST (2023)	Crossformer (2023)	TimesNet (2023)	SCINet (2022a)	DLinear (2023)	FEDformer (2022)	Stationary (2022b)	Autoformer (2021)	Informer (2021)			
Metric	MSE MAE	MSE MAE	MSE MAE	MSE MAE	MSE MAE	MSE MAE	MSE MAE	MSE MAE	MSE MAE	MSE MAE			
Merchant	12	0.058 0.126	0.072 0.155	0.068 0.141	0.088 0.177	0.202 0.310	0.093 0.183	0.277 0.384	0.143 0.243	0.365 0.444	0.489 0.459		
	24	0.066 0.138	0.079	0.164 0.091	0.161	0.103 0.195	0.215 0.323	0.105 0.200	0.268 0.378	0.167 0.270	0.669 0.636	0.507 0.461	
	72	0.079 0.157	0.090 0.180	0.123 0.202	0.089 0.180	0.388 0.431	0.116 0.215	0.281 0.390	0.193 0.300	0.404 0.479	0.523 0.473		
	144	0.086 0.167	0.093 0.185	0.185 0.218	0.091 0.183	0.459 0.477	0.124 0.225	0.359 0.453	0.183 0.294	0.536 0.566	0.543 0.483		
	Avg	0.072 0.147	0.084 0.171	0.117 0.181	0.093 0.184	0.316 0.385	0.110 0.206	0.296 0.401	0.172 0.277	0.494 0.531	0.516 0.469		
Wealth	12	0.189 0.205	0.255	0.250	0.270 0.208	0.275 0.277	0.525 0.451	0.380 0.355	0.553 0.508	0.355 0.332	0.653 0.555	0.837 0.470	
	24	0.254 0.244	0.320 0.291	0.329	0.233	0.300	0.285	0.583 0.479	0.456 0.397	0.567 0.514	0.430 0.377	0.761 0.611	0.905 0.502
	72	0.421	0.327	0.459 0.360	0.484	0.324	0.384 0.326	0.761 0.558	0.555 0.438	0.636 0.548	0.573 0.454	0.857 0.658	1.069 0.573
	144	0.517 0.379	0.541 0.404	0.633 0.388	0.481 0.383	0.770 0.568	0.611 0.459	0.744 0.604	0.637 0.498	0.817 0.627	1.094 0.622		
	Avg	0.345 0.289	0.394 0.326	0.429	0.288	0.360	0.318	0.660 0.514	0.501 0.412	0.625 0.543	0.499 0.415	0.772 0.612	0.976 0.542
Finance	12	0.123 0.170	0.164 0.206	4.630 0.520	0.465 0.291	1.865 0.602	0.321 0.271	1.537 0.538	0.537 0.384	1.651 0.593	7.159 0.911		
	24	0.158 0.197	0.198 0.228	4.987 0.568	0.228 0.297	2.228 0.664	0.464 0.318	1.553 0.547	0.551 0.386	1.671 0.594	7.328 0.915		
	72	0.212 0.240	0.268 0.273	5.631 0.675	0.534 0.310	3.084 0.793	0.986 0.423	1.612 0.554	2.004 0.853	2.054 0.758	7.967 1.004		
	144	0.245 0.257	0.293 0.286	6.083 0.708	0.564 0.333	4.089 0.875	1.287 0.473	1.784 0.636	2.379 0.947	2.114 0.778	7.832 0.995		
	Avg	0.184 0.216	0.231 0.248	5.333 0.618	0.516 0.308	2.817 0.734	0.765 0.372	1.621 0.569	1.368 0.643	1.872 0.681	7.571 0.956		
Terminal	12	0.051 0.127	0.068 0.164	0.055 0.140	0.074 0.169	0.199 0.301	0.096 0.198	0.268 0.379	0.140 0.252	0.386 0.461	0.279 0.365		
	24	0.059 0.139	0.074 0.173	0.065 0.155	0.081 0.178	0.225 0.325	0.105 0.209	0.256 0.370	0.174 0.289	0.708 0.644	0.302 0.378		
	72	0.071 0.160	0.081 0.187	0.077 0.170	0.077 0.178	0.317 0.338	0.109 0.215	0.285 0.396	0.202 0.321	0.510 0.552	0.313 0.377		
	144	0.079 0.171	0.085 0.193	0.085 0.181	0.088 0.192	0.378 0.425	0.113 0.220	0.372 0.468	0.204 0.322	0.468 0.528	0.329 0.393		
	Avg	0.065 0.150	0.077 0.179	0.071 0.162	0.080 0.179	0.280 0.360	0.106 0.210	0.295 0.403	0.180 0.296	0.518 0.547	0.306 0.378		
Payment	12	0.050 0.121	0.065	0.156	0.152 0.145	0.094 0.171	0.164 0.249	0.090 0.180	0.272 0.349	0.129 0.229	0.382 0.437	0.492 0.347	
	24	0.062 0.135	0.077	0.167	0.178 0.165	0.099 0.178	0.216 0.280	0.108 0.196	0.265 0.343	0.157 0.266	0.345 0.412	0.514 0.367	
	72	0.082 0.155	0.094 0.184	0.236 0.193	0.111 0.189	0.360 0.370	0.129 0.209	0.284 0.360	0.183 0.291	0.437 0.471	0.523 0.359		
	144	0.093 0.166	0.101	0.190	0.260 0.214	0.115 0.189	0.410 0.391	0.138 0.215	0.379 0.441	0.194 0.296	0.501 0.518	0.547 0.386	
	Avg	0.072 0.144	0.084 0.174	0.207 0.179	0.105 0.182	0.288 0.322	0.116 0.200	0.300 0.373	0.166 0.271	0.417 0.460	0.519 0.365		
Customer	12	0.065 0.129	0.091	0.160	0.243 0.156	0.123 0.180	0.310 0.326	0.143 0.195	0.309 0.366	0.175 0.243	0.640 0.580	0.828 0.427	
	24	0.078 0.141	0.107 0.173	0.293 0.177	0.130 0.183	0.338 0.344	0.170 0.212	0.313 0.369	0.188 0.264	0.763 0.642	0.860 0.439		
	72	0.108 0.161	0.131 0.190	0.331 0.215	0.149 0.196	0.511 0.408	0.202 0.228	0.330 0.374	0.267 0.324	0.616 0.564	0.904 0.451		
	144	0.126 0.172	0.141 0.195	0.368 0.226	0.166 0.206	0.687 0.461	0.222 0.239	0.450 0.456	0.336 0.373	0.658 0.586	0.903 0.458		
	Avg	0.094 0.150	0.118 0.180	0.309 0.194	0.142 0.191	0.461 0.385	0.184 0.219	0.350 0.391	0.242 0.301	0.669 0.593	0.874 0.444		
1 st Count	28	27	0	0	0	3	2	0	0	0	0	0	